

CSC 580 Artificial Intelligence II, Winter 2025

Name : Omprakash Student ID: 2131025

HW#2 Cliffwalk

1. Install gymnasium and create a Cliff Walk environment

```
%pip install "gymnasium[toy-text]"
%pip install "matplotlib"

import gymnasium as gym
import matplotlib.pyplot as plt
import numpy as np

env = gym.make("CliffWalking-v0", is_slippery=True,
render_mode="rgb_array")
env.reset()

env.render( ) # Get the current frame as an RGB array

array([[[ 91, 150,  97],
        [ 91, 150,  97],
        [ 91, 150,  97],
        ...,
        [ 86, 170,  69],
        [ 86, 170,  69],
        [ 86, 170,  69]],

       [[ 91, 150,  97],
        [ 91, 150,  97],
        [ 91, 150,  97],
        ...,
        [ 86, 170,  69],
        [ 86, 170,  69],
        [ 86, 170,  69]],

       [[ 91, 150,  97],
        [ 91, 150,  97],
        [ 98, 166, 105],
        ...,
        [105, 182,  74],
        [105, 182,  74],
        [ 86, 170,  69]]],

      dtype=uint8)
```

```

    ...,
    [[ 86, 170, 69],
     [ 86, 170, 69],
     [ 93, 188, 74],
     ...,
     [ 91, 150, 97],
     [ 91, 150, 97],
     [ 91, 150, 97]],

    [[ 86, 170, 69],
     [ 86, 170, 69],
     [ 93, 188, 74],
     ...,
     [ 91, 150, 97],
     [ 91, 150, 97],
     [ 91, 150, 97]],

    [[ 86, 170, 69],
     [ 86, 170, 69],
     [ 86, 170, 69],
     ...,
     [ 91, 150, 97],
     [ 91, 150, 97],
     [ 91, 150, 97]]], dtype=uint8)

```

```

#print (env.render()) # textual output if mode == 'ansi' or 'human'
env.render() # gui/rgb output

```

```

array([[[ 91, 150, 97],
         [ 91, 150, 97],
         [ 91, 150, 97],
         ...,
         [ 86, 170, 69],
         [ 86, 170, 69],
         [ 86, 170, 69]],

        [[ 91, 150, 97],
         [ 91, 150, 97],
         [ 91, 150, 97],
         ...,
         [ 86, 170, 69],
         [ 86, 170, 69],
         [ 86, 170, 69]],

        [[ 91, 150, 97],
         [ 91, 150, 97],
         [ 98, 166, 105],

```

```

        ...,
        [105, 182, 74],
        [105, 182, 74],
        [ 86, 170, 69]],
        ...,
        [[ 86, 170, 69],
         [ 86, 170, 69],
         [ 93, 188, 74],
         ...,
         [ 91, 150, 97],
         [ 91, 150, 97],
         [ 91, 150, 97]],
        [[ 86, 170, 69],
         [ 86, 170, 69],
         [ 93, 188, 74],
         ...,
         [ 91, 150, 97],
         [ 91, 150, 97],
         [ 91, 150, 97]],
        [[ 86, 170, 69],
         [ 86, 170, 69],
         [ 86, 170, 69],
         ...,
         [ 91, 150, 97],
         [ 91, 150, 97],
         [ 91, 150, 97]]], dtype=uint8)

```

2. Inspect environment, transition probabilities and rewards

```

nS = env.observation_space.n    # number of states -- 48
nA = env.action_space.n        # number of actions -- four
directions; 0:left, 1:down, 2:right, 3:up
print ("{}", {}".format(nS, nA))

48, 4

```

Transition probabilities, rewards and other info are stored a dictionary **env.P**. You use state index (0-based) to access the info. The tuple for each transition indicates (transition_probability, new_state, reward, terminated).

IMPORTANT: If you get an error "AttributeError: 'OrderEnforcing' object has no attribute 'P'", add a line to unwrap the environment and access P inside. For example,

```

# Probabilities from State 0 (top-left corner).
#env.P[0]

```

```

# Access the underlying environment using env.unwrapped
env_unwrapped = env.unwrapped

# Now you can access the transition probabilities
env_unwrapped.P[0]

{0: [(0.3333333333333333, np.int64(0), -1, False),
      (0.3333333333333333, np.int64(0), -1, False),
      (0.3333333333333333, np.int64(1), -1, False)],
 1: [(0.3333333333333333, np.int64(0), -1, False),
      (0.3333333333333333, np.int64(1), -1, False),
      (0.3333333333333333, np.int64(12), -1, False)],
 2: [(0.3333333333333333, np.int64(1), -1, False),
      (0.3333333333333333, np.int64(12), -1, False),
      (0.3333333333333333, np.int64(0), -1, False)],
 3: [(0.3333333333333333, np.int64(12), -1, False),
      (0.3333333333333333, np.int64(0), -1, False),
      (0.3333333333333333, np.int64(0), -1, False)]}

```

3. Create a random/fixed policy and run the policy once.

```

def generate_random_policy(num_actions, num_states, seed=None):
    """
    A policy is a 1D array of length # of states, where each element
    is a
    number between 0 (inclusive) and # of actions (exclusive) randomly
    chosen.
    If a specific seed is passed, the same numbers are generated,
    while
    if the seed is None, the numbers are unpredictable every time.
    """
    rng = np.random.default_rng(seed)
    return rng.integers(low=0, high=num_actions, size=num_states)

def run(env, pi, printinfo = False):
    """
    Run the policy on the environment and returns the cumulative
    reward.
    :param: env: The environment
    :param: pi: A given policy, represented as a 1D array of length #
    of states.
    :return: Cumulative reward
    """
    s = env.reset()
    if printinfo == True:
        print (f'\n* Episode starting from state {s[0]}') # ensure
        starting from state 36

    s = s[0] # extract the state value/index from the tuple
    done = False # this becomes true when agent reaches the goal

```

```

state (47)
    sum_r = 0
    step_count = 0
    near_fall = 0

    beside_cliff = 0
    cells_around_cliff = [24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,
35,36,47]

    while not done:
        could_be_nearfall = False
        # print('while loop start')
        a = pi[s] # action for the state s, according to the policy
        if s in cells_around_cliff:
            # Check if the chosen action would normally lead to cliff
            # For cells above cliff (24-35), check if action is "down"
(1)            # For cell 36, check if action is "right" (2)
            # For cell 47, check if action is "left" (0)
            beside_cliff += 1
            if (s in range(24, 36) and a == 1) or (s == 36 and a == 2)
or (s == 47 and a == 0):
                could_be_nearfall = True

        step_count += 1
        s, r, done, info, p = env.step(a) # take the action

        sum_r += r # accumulate reward

        #Here my thought process was
        #if the agent is near the cliff and the reward is not -100,
        then it is a near fall and missed the fall
        #I am not sure if this is the correct way to check for near
        fall
        if could_be_nearfall and r != -100:
            near_fall += 1

        ### uncomment below to see the information for each step
        #print (f'next_state={s}, reward={r}, done={done},
        info={info}, p={p}')

        # prints info in text if render_mode is 'ansi' or no output if
        'human',
        # or graphical output if 'rgb_array' AND if the code is run
        from command line.
        #env.render()
        return sum_r, step_count, beside_cliff/step_count,
        near_fall/step_count

```

Run the given policy just once and observe what's returned (the total return).

```
policy = generate_random_policy(nA, nS, 17) # third parameter is the
random seed
print ("*** Policy ***\n{}".format(policy.reshape((4, 12))))

# Do just one run
result = run(env, policy)
# Print the total rewards/return
print (f' ==> Total return: {result}')
```

*** Policy ***

```
[[2 3 0 0 1 2 3 1 0 0 1 1]
 [3 1 2 2 0 2 0 0 2 1 1 2]
 [0 0 2 3 2 3 1 0 2 2 2 2]
 [1 0 2 0 1 0 2 1 1 2 0 1]]
==> Total return: (-50824, 8353, 0.4327786424039267, 0.0)
```

```
# return sum_r, step_count, beside_cliff/step_count,
near_fall/step_count
from tqdm import tqdm
def procedure(policy, num_episodes):
    """
    Run the policy for a number of episodes and return the average
    cumulative reward.
    :param: policy: A given policy, represented as a 1D array of
    length # of states.
    """
    print ("Running the policy for {}
episodes...".format(num_episodes))
    total_reward = []
    total_steps = []
    total_near_fall = []
    total_beside_cliff = []
    for i in tqdm(range(num_episodes), position=0, leave=True):
        r, steps, beside_cliff, near_fall_count = run(env, policy)
        total_beside_cliff.append(beside_cliff)
        total_reward.append(r)
        total_steps.append(steps)
        total_near_fall.append(near_fall_count)
    return total_reward, total_steps, total_near_fall,
total_beside_cliff

import matplotlib.pyplot as plt
def output(total_reward, total_steps, total_near_fall,
total_beside_cliff, seed):
    """
    Output the average cumulative reward and the standard deviation.
    """
```

```

print(f' ==> Output for seed: {seed}')
print (f' ==> Average reward: {np.mean(total_reward)}')
print (f' ==> Average steps: {np.mean(total_steps)}')
print (f' ==> Average near fall: {np.mean(total_near_fall)}')
print (f' ==> Average beside cliff:
{np.mean(total_beside_cliff)}')

print (f' ==> Standard deviation of reward:
{np.std(total_reward)}')
print (f' ==> Standard deviation of steps: {np.std(total_steps)}')
print (f' ==> Standard deviation of near fall:
{np.std(total_near_fall)}')
print (f' ==> Standard deviation of beside cliff:
{np.std(total_beside_cliff)}')


fig, axs = plt.subplots(1, 4, figsize=(12, 6))

# Steps
axs[0].hist(total_steps, bins=20, color='orange',
edgecolor='black')
axs[0].set_title("Steps for seed: {}".format(seed))
axs[0].set_yscale('linear') # Set y-scale to linear

# Ratio of Beside cliff
axs[1].hist(total_beside_cliff, bins=20, color='brown',
edgecolor='black')
axs[1].set_title("Ratio of Around cliff for seed:
{}".format(seed))

# Rewards
axs[2].hist(total_reward, bins=20, color='green',
edgecolor='black')
axs[2].set_title("Rewards for seed: {}".format(seed))

# Ratio of Near-falls
axs[3].hist(total_near_fall, bins=20, color='cyan',
edgecolor='black')
axs[3].set_title("Ratio of Near-falls for seed: {}".format(seed))

plt.tight_layout() # Adjust layout
plt.show()

# Do 100 runs and print the average return
n = 100

```


```
policy1 = generate_random_policy(nA, nS, 17) # third parameter is the
random seed
```

```
for policy in [policy1]:
    total_reward, total_steps, total_near_fall, total_beside_cliff =
    procedure(policy, n)
    output(total_reward, total_steps,
    total_near_fall, total_beside_cliff, seed=17)
```

Running the policy for 100 episodes...

```
100%|██████████| 100/100 [00:01<00:00, 72.94it/s]
```

```
==> Output for seed: 17
==> Average reward: -20748.47
==> Average steps: 3127.46
==> Average near fall: 5.905898453590458e-05
==> Average beside cliff: 0.4136403416570851
==> Standard deviation of reward: 20842.677662169513
==> Standard deviation of steps: 3045.672285128523
==> Standard deviation of near fall: 0.00020718032186367992
==> Standard deviation of beside cliff: 0.049851714076945775
```



```
policy2 = generate_random_policy(nA, nS, 12) # third parameter is the
random seed
policy3 = generate_random_policy(nA, nS, 23) # third parameter is the
random seed
policy4 = generate_random_policy(nA, nS, 25) # third parameter is the
random seed
policies = [policy2, policy3, policy4]
```



```

seeds = [19, 21, 23]
for policy, seed in zip(policies, seeds):
    total_reward, total_steps, total_near_fall, total_beside_cliff =
    procedure(policy, n)
    output(total_reward, total_steps,
    total_near_fall, total_beside_cliff, seed)

```


Running the policy for 100 episodes...

100%|██████████| 100/100 [00:01<00:00, 83.74it/s]

```

==> Output for seed: 19
==> Average reward: -25970.7
==> Average steps: 2706.69
==> Average near fall: 0.14637418438253977
==> Average beside cliff: 0.4630669640132981
==> Standard deviation of reward: 27633.31632558785
==> Standard deviation of steps: 2809.723810964344
==> Standard deviation of near fall: 0.019815597099205644
==> Standard deviation of beside cliff: 0.05910016673836438

```



Running the policy for 100 episodes...


100%|██████████| 100/100 [00:00<00:00, 349.19it/s]

```

==> Output for seed: 21
==> Average reward: -3166.11
==> Average steps: 639.63
==> Average near fall: 0.0658426764757201
==> Average beside cliff: 0.21353544885531203
==> Standard deviation of reward: 3125.727307027918

```


```
==> Standard deviation of steps: 581.7097498753137
==> Standard deviation of near fall: 0.020841470627897433
==> Standard deviation of beside cliff: 0.057665078024030626
```



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:01<00:00, 50.42it/s]

```
==> Output for seed: 23
==> Average reward: -43920.51
==> Average steps: 4552.17
==> Average near fall: 0.12307805207407395
==> Average beside cliff: 0.37861339218735807
==> Standard deviation of reward: 38895.25873020901
==> Standard deviation of steps: 4004.7783710337826
==> Standard deviation of near fall: 0.022057890512962967
==> Standard deviation of beside cliff: 0.06594869929257952
```



Policy Iteration

```
import numpy as np
def getInitialPolicy(seed):
    V = np.array([0.1 if i != 47 else 0 for i in range(48)])
    theta = 1e-6 # Threshold for convergence
    gamma = 0.8 # Discount factor
    print("Initial Value Function V(s):")
    print(V.reshape(4, 12))
    policy = generate_random_policy(nA, nS, seed) # third parameter is
the random seed
    print('Initial Policy:')
    print(policy.reshape(4, 12))

    return V, policy, seed

def policy_evaluation(env, V, policy, gamma, theta):
    while True:
        delta = 0
        # Loop for each state
        for s in range(nS):
            v = V[s]
            V[s] = sum(
                prob * (reward + (0 if done else gamma *
V[next_state]))
                for prob, next_state, reward, done in env.P[s]
                [policy[s]]
            )
            delta = max(delta, abs(v - V[s]))
        if delta < theta:
```

```

        break
    return V

def policy_improvement(env, V, policy, gamma):
    # Step 3: Policy Improvement
    policy_stable = True
    for s in range(len(policy)):
        old_action = policy[s]

        action_values = []
        for a in range(nA):
            action_value = sum(
                prob * (reward + (0 if done else gamma *
V[next_state]))
                for prob, next_state, reward, done in env.P[s][a]
            )
            action_values.append(action_value)

        policy[s] = max(range(len(action_values)), key=lambda a:
action_values[a])
        if old_action != policy[s]:
            policy_stable = False
            # print(f"State {s:2d} | {old_action} → {policy[s]} | Q:
{[f'{q:.3f}' for q in action_values]}")
        return policy, policy_stable

def policy_iteration(seed):
    env.reset()
    V, policy, seed = getInitialPolicy(seed)
    gamma = 0.8 # Discount factor
    theta = 1e-6 # Threshold for convergence
    iter = 0
    while True:
        V = policy_evaluation(env.unwrapped, V, policy, gamma, theta)
        policy, policy_stable = policy_improvement(env.unwrapped, V,
policy, gamma)
        iter += 1
        if policy_stable:
            break
    return V, policy, seed, iter

for seed in [17, 19, 21, 23, 25]:
    optimal_V, optimal_policy, seed, iter = policy_iteration(seed)
    total_reward, total_steps, total_near_fall, total_beside_cliff =
procedure(optimal_policy, n)
    print(f'==> Seed: {seed}, is converged in {iter} iterations')
    print("Optimal Value Function V(s): \n", optimal_V.reshape(4, 12),
"\n")
    print("Optimal Policy: \n", optimal_policy.reshape(4, 12), "\n")

```

```
output(total_reward, total_steps,
total_near_fall,total_beside_cliff,seed)
```

Initial Value Function V(s):

```
[[0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0. ]]
```

Initial Policy:

```
[[2 3 0 0 1 2 3 1 0 0 1 1]
 [3 1 2 2 0 2 0 0 2 1 1 2]
 [0 0 2 3 2 3 1 0 2 2 2 2]
 [1 0 2 0 1 0 2 1 1 2 0 1]]
```

Running the policy for 100 episodes...

```
100%|██████████| 100/100 [00:00<00:00, 2352.94it/s]
```

==> Seed: 17, is converged in 17 iterations

Optimal Value Function V(s):

```
[[ -4.9986267  -4.99759669  -4.99553287  -4.99168868  -4.98457607
  -4.97152048  -4.94786452  -4.90594181  -4.83449661  -4.72121416
  -4.56594567  -4.4181314 ]
 [ -4.99872172  -4.99785799  -4.99602668  -4.99256779  -4.9860637
  -4.9738168   -4.9506856   -4.90684335  -4.8236515   -4.66739328
  -4.38821919  -3.98172995]
 [ -4.99893303  -4.99834411  -4.99699937  -4.99437684  -4.98934599
  -4.97960692  -4.96036314  -4.92106927  -4.83680325  -4.64329143
  -4.15814634  -2.81153816]
 [ -4.99939032 -37.99923326 -70.99887467 -70.99817533 -70.99683377
  -70.99423668 -70.98910501 -70.97862664 -70.9561557  -70.90455255
  -36.44200978  -1.74974351]]
```

Optimal Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

==> Output for seed: 17

==> Average reward: -65.41

==> Average steps: 65.41

==> Average near fall: 0.048114485120904577


==> Average beside cliff: 0.35647496564551434

==> Standard deviation of reward: 27.37849338440667

==> Standard deviation of steps: 27.37849338440667

==> Standard deviation of near fall: 0.029948945523365248

==> Standard deviation of beside cliff: 0.10797118128418773



Initial Value Function V(s):

```
[[0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0. ]]]
```

Initial Policy:

```
[[2 1 1 3 1 1 3 0 1 1 0 2]
 [3 3 1 2 3 1 1 3 0 3 2 1]
 [3 1 1 2 3 2 2 2 3 1 2 1]
 [2 2 1 3 0 0 2 0 2 3 0 1]]
```

Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2823.69it/s]

==> Seed: 19, is converged in 4 iterations

Optimal Value Function V(s):

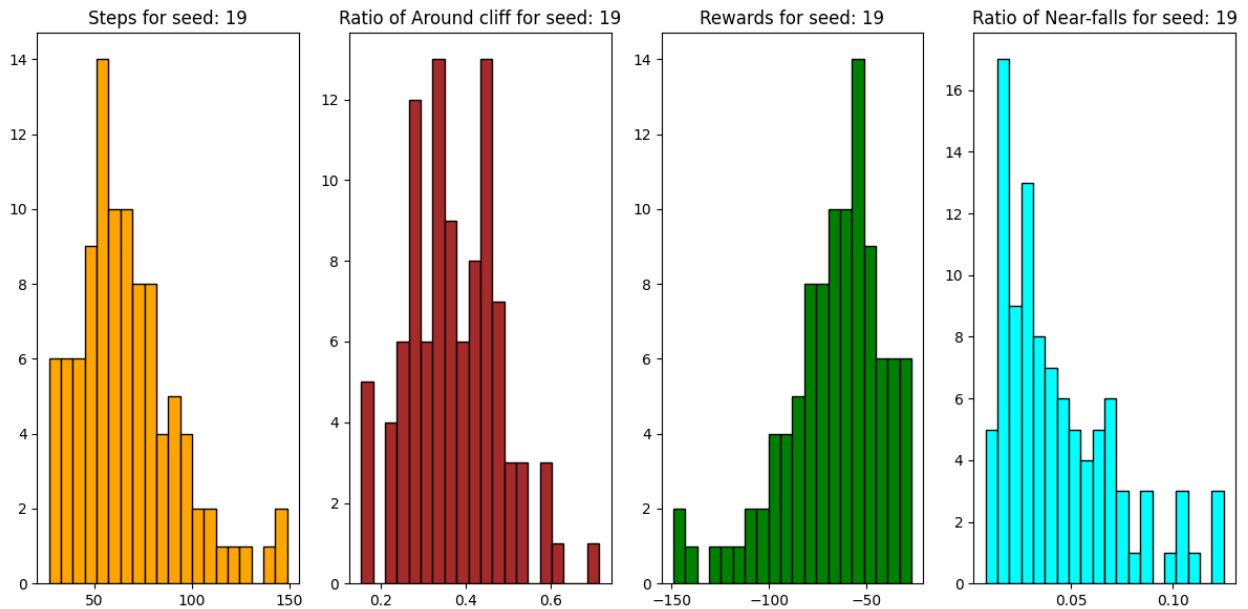
```
[[ -4.99862897  -4.99759853  -4.99553461  -4.99169024  -4.98457739
  -4.97152156  -4.94786536  -4.90594245  -4.8344971  -4.72121454
  -4.56594597  -4.41813166]
 [ -4.99872343  -4.99785935  -4.99602792  -4.99256889  -4.98606462
  -4.97381755  -4.95068619  -4.9068438  -4.82365184  -4.66739352
  -4.38821936  -3.98173006]
 [ -4.99893434  -4.99834517  -4.9970003  -4.99437765  -4.98934668
  -4.97960748  -4.96036359  -4.92106961  -4.83680351  -4.64329161
  -4.15814646  -2.81153821]
 [ -4.99939142 -37.99923414 -70.9988755  -70.99817613 -70.99683454
  -70.99423742 -70.98910572 -70.97862732 -70.95615636 -70.90455319
  -36.4420101  -1.74974352]]
```

Optimal Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]]
```

```
[0 1 1 1 1 1 1 1 1 1 1 2]
[0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 0 0 0 0 0 0 0 0 0 1]]
```

```
==> Output for seed: 19
==> Average reward: -67.19
==> Average steps: 67.19
==> Average near fall: 0.04423482257792665
==> Average beside cliff: 0.37174298713181125
==> Standard deviation of reward: 25.50674224592392
==> Standard deviation of steps: 25.50674224592392
==> Standard deviation of near fall: 0.028083632648032365
==> Standard deviation of beside cliff: 0.10720112749193558
```



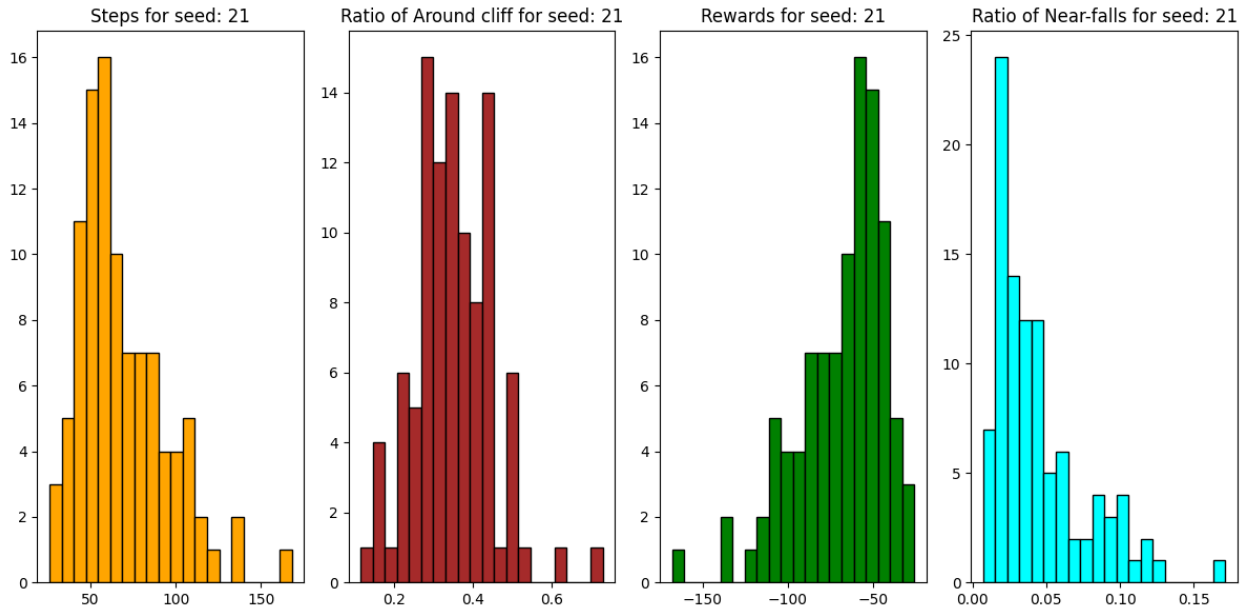
```
Initial Value Function V(s):
[[0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0. ]]
Initial Policy:
[[1 3 1 2 1 2 1 0 1 2 2 3]
 [0 1 3 0 2 3 3 2 3 0 3 2]
 [2 3 3 0 3 3 1 2 3 0 1 0]
 [1 3 1 1 3 1 0 2 3 1 3 0]]
Running the policy for 100 episodes...
```

```
100%|██████████| 100/100 [00:00<00:00, 2174.34it/s]
```

```
==> Seed: 21, is converged in 8 iterations
Optimal Value Function V(s):
[[ -4.9986276   -4.99759733  -4.99553343  -4.99168916  -4.98457646
  -4.97152079  -4.94786475  -4.90594198  -4.83449674  -4.72121427
  -4.56594575  -4.41813147]
 [ -4.99872237  -4.99785846  -4.99602708  -4.99256813  -4.98606397
  -4.97381702  -4.95068577  -4.90684348  -4.8236516   -4.66739334
  -4.38821924  -3.98172998]
 [ -4.99893353  -4.99834448  -4.99699967  -4.99437709  -4.98934619
  -4.97960708  -4.96036327  -4.92106936  -4.83680333  -4.64329148
  -4.15814637  -2.81153818]
 [ -4.99939078 -37.99923361 -70.998875   -70.99817564 -70.99683407
  -70.99423697 -70.98910529 -70.97862691 -70.95615597 -70.90455281
  -36.44200991  -1.74974351]]
```

```
Optimal Policy:
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

```
==> Output for seed: 21
==> Average reward: -68.54
==> Average steps: 68.54
==> Average near fall: 0.04444377812547879
==> Average beside cliff: 0.35029821153256335
==> Standard deviation of reward: 25.854369069849685
==> Standard deviation of steps: 25.854369069849685
==> Standard deviation of near fall: 0.03131589652718283
==> Standard deviation of beside cliff: 0.09902893720368824
```

Initial Value Function $V(s)$:

```
[[0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0. ]]
```

Initial Policy:

```
[[0 2 1 2 1 0 2 0 0 2 2 3]
 [1 0 0 0 0 2 3 1 2 1 2 1]
 [1 0 0 1 0 1 0 1 0 2 1 2]
 [0 2 0 2 1 0 1 2 1 2 3 3]]
```

Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2609.75it/s]

==> Seed: 23, is converged in 7 iterations

Optimal Value Function $V(s)$:


```
[[ -4.9986276  -4.99759733 -4.99553343 -4.99168916 -4.98457646
  -4.97152079 -4.94786475 -4.90594198 -4.83449674 -4.72121427
  -4.56594575 -4.41813147]
 [ -4.99872237 -4.99785846 -4.99602708 -4.99256813 -4.98606397
  -4.97381702 -4.95068577 -4.90684348 -4.8236516  -4.66739334
  -4.38821924 -3.98172998]
 [ -4.99893353 -4.99834448 -4.99699967 -4.99437709 -4.98934619
  -4.97960708 -4.96036327 -4.92106936 -4.83680332 -4.64329148
  -4.15814637 -2.81153818]
 [ -4.99939078 -37.99923361 -70.998875  -70.99817564 -70.99683407
  -70.99423697 -70.98910529 -70.97862691 -70.95615597 -70.90455281
  -36.44200991 -1.74974351]]
```

Optimal Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]]
```

```
[0 1 1 1 1 1 1 1 1 1 1 2]
[0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 0 0 0 0 0 0 0 0 0 1]]
```

```
==> Output for seed: 23
==> Average reward: -66.6
==> Average steps: 66.6
==> Average near fall: 0.045284838491584115
==> Average beside cliff: 0.37952300813775947
==> Standard deviation of reward: 22.78859363804621
==> Standard deviation of steps: 22.78859363804621
==> Standard deviation of near fall: 0.03412954741939584
==> Standard deviation of beside cliff: 0.13042366686020418
```



Initial Value Function V(s):

```
[[0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0. ]]
```

Initial Policy:

```
[[2 0 3 0 0 0 2 1 0 0 2 0]
 [1 3 2 3 2 0 0 0 0 3 0 3]
 [1 1 0 3 1 0 2 2 1 2 2 3]
 [1 1 3 1 1 1 2 0 1 2 3 2]]
```

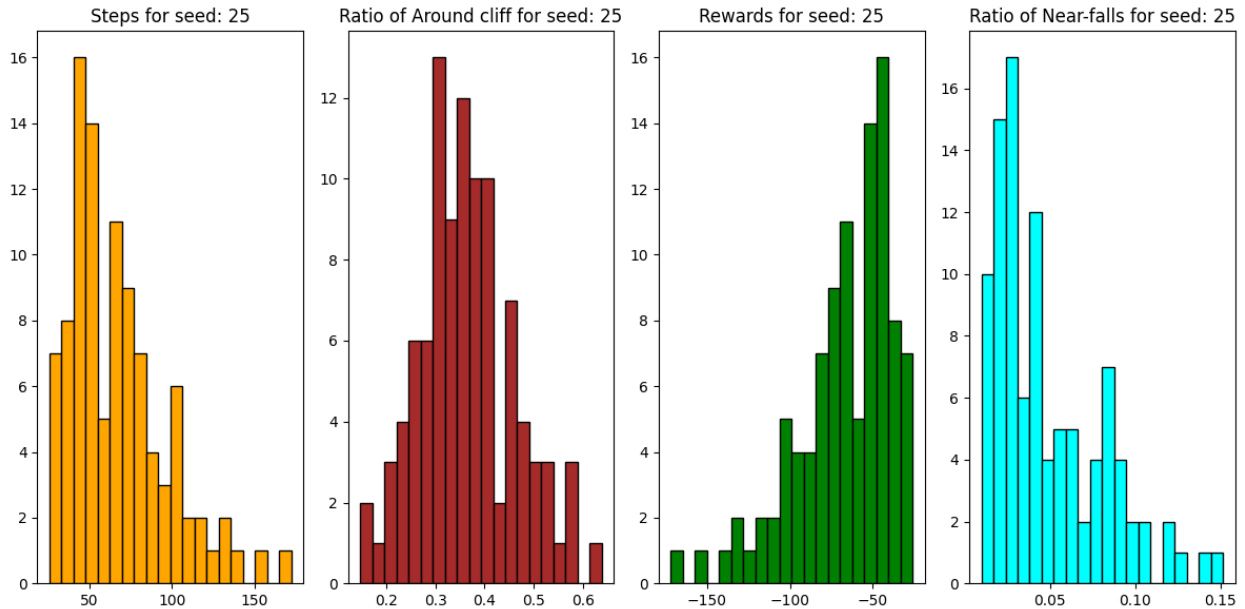
Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2783.27it/s]

```
==> Seed: 25, is converged in 8 iterations
Optimal Value Function V(s):
[[ -4.9986276   -4.99759733  -4.99553343  -4.99168916  -4.98457646
   -4.97152079  -4.94786475  -4.90594198  -4.83449674  -4.72121427
   -4.56594575  -4.41813147]
 [ -4.99872237  -4.99785846  -4.99602708  -4.99256813  -4.98606397
   -4.97381702  -4.95068577  -4.90684348  -4.8236516   -4.66739334
   -4.38821924  -3.98172998]
 [ -4.99893353  -4.99834448  -4.99699967  -4.99437709  -4.98934619
   -4.97960708  -4.96036327  -4.92106936  -4.83680333  -4.64329148
   -4.15814637  -2.81153818]
 [ -4.99939078 -37.99923361 -70.998875   -70.99817564 -70.99683407
  -70.99423697 -70.98910529 -70.97862691 -70.95615597 -70.90455281
  -36.44200991  -1.74974351]]
```

```
Optimal Policy:
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

```
==> Output for seed: 25
==> Average reward: -67.27
==> Average steps: 67.27
==> Average near fall: 0.04852925955657054
==> Average beside cliff: 0.36290444000994704
==> Standard deviation of reward: 29.009948293645753
==> Standard deviation of steps: 29.009948293645753
==> Standard deviation of near fall: 0.03187937444039755
==> Standard deviation of beside cliff: 0.09729814367396875
```



Try with various $V(s)$

```
def policy_iteration_with_init_v(seed, init_v_value):
    env.reset()
    # Initialize V with the specified value except for terminal state
    V = np.array([init_v_value if i != 47 else 0 for i in range(48)])
    policy = generate_random_policy(nA, nS, seed)

    print(f"\nInitial setup for seed {seed} with
V(s)={init_v_value}:")
    print("Initial Value Function V(s):")
    print(V.reshape(4, 12))
    print('Initial Policy:')
    print(policy.reshape(4, 12))

    gamma = 0.8
    theta = 1e-6
    iter = 0

    while True:
        V = policy_evaluation(env.unwrapped, V, policy, gamma, theta)
        policy, policy_stable = policy_improvement(env.unwrapped, V,
policy, gamma)
        iter += 1
        if policy_stable:
            break

    return V, policy, iter

# Test different initial V(s) values
init_values = [0.1, 0.5, 1.0]
```

```

seeds = [17]

for seed in seeds:
    print(f"\n=== Testing seed {seed} with different initial V(s) values ===")
    for init_v in init_values:
        optimal_V, optimal_policy, iter =
policy_iteration_with_init_v(seed, init_v)
        total_reward, total_steps, total_near_fall, total_beside_cliff
= procedure(optimal_policy, n)
        print(f'\nResults for initial V(s)={init_v}:')
        print(f'Converged in {iter} iterations')
        print("Optimal Value Function V(s): \n", optimal_V.reshape(4,
12), "\n")
        print("Optimal Policy: \n", optimal_policy.reshape(4, 12), "\n")
        output(total_reward, total_steps, total_near_fall,
total_beside_cliff, seed)

```

=== Testing seed 17 with different initial V(s) values ===

Initial setup for seed 17 with V(s)=0.1:

Initial Value Function V(s):

```

[[0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1]
 [0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0. ]]]

```

Initial Policy:

```

[[2 3 0 0 1 2 3 1 0 0 1 1]
 [3 1 2 2 0 2 0 0 2 1 1 2]
 [0 0 2 3 2 3 1 0 2 2 2 2]
 [1 0 2 0 1 0 2 1 1 2 0 1]]

```

Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2608.85it/s]

Results for initial V(s)=0.1:

Converged in 17 iterations

Optimal Value Function V(s):

```

[[ -4.9986267  -4.99759669  -4.99553287  -4.99168868  -4.98457607
  -4.97152048  -4.94786452  -4.90594181  -4.83449661  -4.72121416
  -4.56594567  -4.4181314 ]
 [ -4.99872172  -4.99785799  -4.99602668  -4.99256779  -4.9860637
  -4.9738168   -4.9506856   -4.90684335  -4.8236515   -4.66739328
  -4.38821919  -3.98172995]
 [ -4.99893303  -4.99834411  -4.99699937  -4.99437684  -4.98934599
  -4.97960692  -4.96036314  -4.92106927  -4.83680325  -4.64329143
  -4.15814634  -2.81153816]]

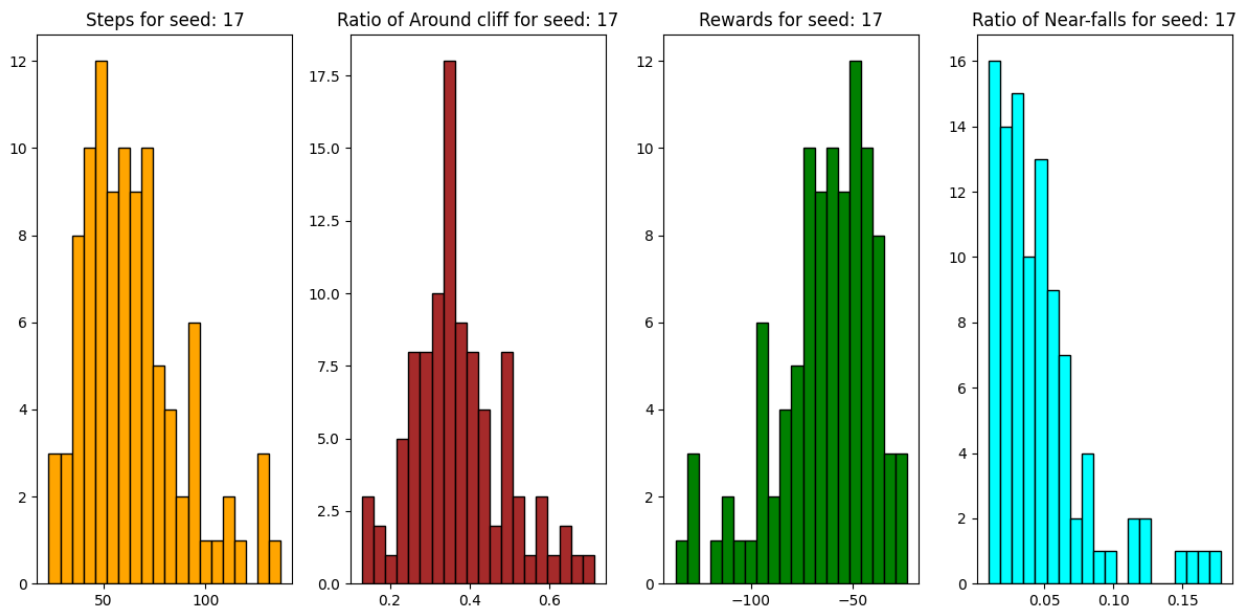
```

```
[ -4.99939032 -37.99923326 -70.99887467 -70.99817533 -70.99683377
-70.99423668 -70.98910501 -70.97862664 -70.9561557 -70.90455255
-36.44200978 -1.74974351]]
```

Optimal Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
[0 1 1 1 1 1 1 1 1 1 1 2]
[0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 0 0 0 0 0 0 0 0 0 1]]
```

```
==> Output for seed: 17
==> Average reward: -63.94
==> Average steps: 63.94
==> Average near fall: 0.04681296298984998
==> Average beside cliff: 0.3739095837959356
==> Standard deviation of reward: 24.427369895262977
==> Standard deviation of steps: 24.427369895262977
==> Standard deviation of near fall: 0.03448101391562455
==> Standard deviation of beside cliff: 0.11532566697113665
```



Initial setup for seed 17 with $V(s)=0.5$:

Initial Value Function $V(s)$:

```
[[0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5]
[0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5]
[0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5]
[0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0. ]]
```

Initial Policy:

```
[[2 3 0 0 1 2 3 1 0 0 1 1]]
```

```
[3 1 2 2 0 2 0 0 2 1 1 2]
[0 0 2 3 2 3 1 0 2 2 2 2]
[1 0 2 0 1 0 2 1 1 2 0 1]]
```

Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2551.99it/s]

Results for initial $V(s)=0.5$:

Converged in 17 iterations

Optimal Value Function $V(s)$:

```
[[ -4.9986267  -4.99759669 -4.99553287  -4.99168868  -4.98457607
  -4.97152048  -4.94786452  -4.90594181  -4.83449661  -4.72121416
  -4.56594567  -4.4181314 ]
 [ -4.99872172  -4.99785799  -4.99602668  -4.99256779  -4.9860637
  -4.9738168   -4.9506856   -4.90684335  -4.8236515   -4.66739328
  -4.38821919  -3.98172995]
 [ -4.99893303  -4.99834411  -4.99699937  -4.99437684  -4.98934599
  -4.97960692  -4.96036314  -4.92106927  -4.83680325  -4.64329143
  -4.15814634  -2.81153816]
 [ -4.99939032 -37.99923326 -70.99887467 -70.99817533 -70.99683377
  -70.99423668 -70.98910501 -70.97862664 -70.9561557  -70.90455255
  -36.44200978  -1.74974351]]
```

Optimal Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

==> Output for seed: 17

==> Average reward: -68.64

==> Average steps: 68.64

==> Average near fall: 0.04595408376246253


==> Average beside cliff: 0.36735147081127695

==> Standard deviation of reward: 30.24285700789527

==> Standard deviation of steps: 30.24285700789527

==> Standard deviation of near fall: 0.03201574519793338

==> Standard deviation of beside cliff: 0.10777689667480266



Initial setup for seed 17 with $V(s)=1.0$:

Initial Value Function $V(s)$:

```
[[1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0.]]
```

Initial Policy:

```
[[2 3 0 0 1 2 3 1 0 0 1 1]
 [3 1 2 2 0 2 0 0 2 1 1 2]
 [0 0 2 3 2 3 1 0 2 2 2 2]
 [1 0 2 0 1 0 2 1 1 2 0 1]]
```

Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2656.32it/s]

Results for initial $V(s)=1.0$:

Converged in 17 iterations

Optimal Value Function $V(s)$:

```
[[ -4.9986267  -4.99759669 -4.99553287  -4.99168868  -4.98457607
  -4.97152048  -4.94786452  -4.90594181  -4.83449661  -4.72121416
  -4.56594567  -4.4181314 ]
 [ -4.99872172  -4.99785799  -4.99602668  -4.99256779  -4.9860637
  -4.9738168   -4.9506856   -4.90684335  -4.8236515   -4.66739328
  -4.38821919  -3.98172995]
 [ -4.99893303  -4.99834411  -4.99699937  -4.99437684  -4.98934599
  -4.97960692  -4.96036314  -4.92106927  -4.83680325  -4.64329143
  -4.15814634  -2.81153816]
 [ -4.99939032 -37.99923326 -70.99887467 -70.99817533 -70.99683377
  -70.99423668 -70.98910501 -70.97862664 -70.9561557  -70.90455255]
```




```
-36.44200978 -1.74974351]]
```

Optimal Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

```
==> Output for seed: 17
==> Average reward: -61.45
==> Average steps: 61.45
==> Average near fall: 0.047094506516299706
==> Average beside cliff: 0.35372322164570236
==> Standard deviation of reward: 21.31730517678067
==> Standard deviation of steps: 21.31730517678067
==> Standard deviation of near fall: 0.028863020349988402
==> Standard deviation of beside cliff: 0.1106394129379537
```



Test different theta values

```
# Test different theta values
theta_values = [1e-2, 1e-4, 1e-6, 1e-8]
seed = 17 # Use a fixed seed for comparison

results = []
for theta in theta_values:
    env.reset()
    V = np.array([0.1 if i != 47 else 0 for i in range(48)])
    policy = generate_random_policy(nA, nS, seed)
```

```
gamma = 0.8

# Modified policy iteration with different theta
iter = 0
while True:
    V = policy_evaluation(env.unwrapped, V, policy, gamma, theta)
    policy, policy_stable = policy_improvement(env.unwrapped, V,
policy, gamma)
    iter += 1
    if policy_stable:
        break

# Run the policy to evaluate performance
total_reward, total_steps, total_near_fall, total_beside_cliff =
procedure(policy, n)

print(f"\nResults for θ={theta}:")
print(f'Converged in {iter} iterations')
print("Final Value Function V(s):")
print(V.reshape(4, 12))
print("\nFinal Policy:")
print(policy.reshape(4, 12))

output(total_reward, total_steps, total_near_fall,
total_beside_cliff, seed)

# Store results for comparison
results.append({
    'theta': theta,
    'iterations': iter,
    'avg_reward': np.mean(total_reward),
    'avg_steps': np.mean(total_steps)
})

# Print comparison summary
print("\nComparison Summary:")
print("θ\tIterations\tAvg Reward\tAvg Steps")
print("-" * 50)
for r in results:
    print(f"{r['theta']:.0e}\t{r['iterations']}\t\
t{r['avg_reward']:.2f}\t{r['avg_steps']:.2f}")
```

Running the policy for 100 episodes...

Progress	Total Rewards	Avg Reward	Total Steps	Avg Steps	Near Fall Count	Beside Cliff Count
100% ██████████	100/100	[00:00<00:00]	2313.31	lit/s]		

Results for θ=0.01:
 Converged in 7 iterations
 Final Value Function V(s):

```
[ [-4.99974879 -4.99917669 -4.99785683 -4.99495761 -4.98897385
-4.97739295 -4.95562237 -4.91497136 -4.84303408 -4.72872373
-4.57427215 -4.42704259]
[ -4.99972277 -4.99914467 -4.99786979 -4.99511843 -4.98950763
-4.97846512 -4.95659452 -4.91323948 -4.82962714 -4.67339795
-4.39450166 -3.98706661]
[ -4.99972886 -4.99931131 -4.99832532 -4.99619625 -4.99182603
-4.9829403 -4.9644863 -4.92548895 -4.84105703 -4.64758973
-4.16212985 -2.81412324]
[ -4.99991333 -37.99977012 -70.99950719 -70.99893944 -70.99777405
-70.99540452 -70.99048346 -70.98008416 -70.95756898 -70.9059777
-36.44321151 -1.75043286]]
```

Final Policy:

```
[[1 1 1 1 1 1 1 1 1 1 1 1]
[2 1 1 1 1 1 1 1 1 1 1 2]
[0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 3 0 0 0 0 0 0 0 0 1]]
```

==> Output for seed: 17

==> Average reward: -65.92

==> Average steps: 65.92

==> Average near fall: 0.03927888569702545


==> Average beside cliff: 0.40331767481258796

==> Standard deviation of reward: 30.288175910741142

==> Standard deviation of steps: 30.288175910741142

==> Standard deviation of near fall: 0.02637532091045519

==> Standard deviation of beside cliff: 0.11894505068064007



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2602.14it/s]

Results for $\theta=0.0001$:

Converged in 7 iterations

Final Value Function $V(s)$:

```
[[ -4.99864224 -4.99762088 -4.99557299 -4.99174982 -4.98465864
  -4.97161808 -4.94796583 -4.90603573 -4.83457653 -4.72127894
  -4.56599818 -4.41817685]
 [ -4.9987356 -4.99787862 -4.99605947 -4.99261549 -4.98612526
  -4.97388679 -4.95075613 -4.90690731 -4.8237047 -4.66743455
  -4.38824919 -3.98174971]
 [ -4.99894439 -4.99835968 -4.99702288 -4.99441025 -4.98938865
  -4.97965541 -4.96041244 -4.92111456 -4.83684123 -4.64332047
  -4.15816575 -2.81154709]
 [ -4.99939749 -37.99924124 -70.99888476 -70.99818806 -70.99684897
  -70.99425344 -70.98912198 -70.97864254 -70.95616966 -70.90456412
  -36.44201686 -1.74974589]]
```

Final Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

==> Output for seed: 17

==> Average reward: -68.81

==> Average steps: 68.81

==> Average near fall: 0.048958052577079025


==> Average beside cliff: 0.37586691545572265

==> Standard deviation of reward: 26.095093408531806

==> Standard deviation of steps: 26.095093408531806

==> Standard deviation of near fall: 0.033862313919982345

==> Standard deviation of beside cliff: 0.09927549521910899



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2609.68it/s]

Results for $\theta=1e-06$:

Converged in 17 iterations

Final Value Function V(s):

[-4.9986267	-4.99759669	-4.99553287	-4.99168868	-4.98457607
-4.97152048	-4.94786452	-4.90594181	-4.83449661	-4.72121416
-4.56594567	-4.4181314]			
[-4.99872172	-4.99785799	-4.99602668	-4.99256779	-4.9860637
-4.9738168	-4.9506856	-4.90684335	-4.8236515	-4.66739328
-4.38821919	-3.98172995]			
[-4.99893303	-4.99834411	-4.99699937	-4.99437684	-4.98934599
-4.97960692	-4.96036314	-4.92106927	-4.83680325	-4.64329143
-4.15814634	-2.81153816]			
[-4.99939032	-37.99923326	-70.99887467	-70.99817533	-70.99683377
-70.99423668	-70.98910501	-70.97862664	-70.9561557	-70.90455255
-36.44200978	-1.74974351]]			

Final Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 0 1]]
```

==> Output for seed: 17

==> Average reward: -65.53

==> Average steps: 65.53


==> Average near fall: 0.04547643340822742

==> Average beside cliff: 0.38316216248068635


```

[0 0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 0 0 0 0 0 0 0 0 0 0 1]]
==> Output for seed: 17
==> Average reward: -66.11
==> Average steps: 66.11
==> Average near fall: 0.05132612617895859
==> Average beside cliff: 0.37064503270312116
==> Standard deviation of reward: 20.874814969239846
==> Standard deviation of steps: 20.874814969239846
==> Standard deviation of near fall: 0.03898635942745567
==> Standard deviation of beside cliff: 0.10206355667329345

```



Comparison Summary:

θ	Iterations	Avg Reward	Avg Steps
----------	------------	------------	-----------

1e-02	7	-65.92	65.92
1e-04	7	-68.81	68.81
1e-06	17	-65.53	65.53
1e-08	27	-66.11	66.11

Try different gamma

```

# Test different gamma values
gamma_values = [0.1, 0.5, 0.8, 0.9, 0.99]
seed = 17 # Use a fixed seed for comparison

results = []

```

```
for gamma in gamma_values:
    env.reset()
    V = np.array([0.1 if i != 47 else 0 for i in range(48)])
    policy = generate_random_policy(nA, nS, seed)
    theta = 1e-6 # Fixed theta value

# Policy iteration with different gamma
iter = 0
while True:
    V = policy_evaluation(env.unwrapped, V, policy, gamma, theta)
    policy, policy_stable = policy_improvement(env.unwrapped, V,
policy, gamma)
    iter += 1
    if policy_stable:
        break

# Run the policy to evaluate performance
total_reward, total_steps, total_near_fall, total_beside_cliff =
procedure(policy, n)

print(f"\nResults for γ={gamma}:")
print(f'Converged in {iter} iterations')
print("Final Value Function V(s):")
print(V.reshape(4, 12))
print("\nFinal Policy:")
print(policy.reshape(4, 12))

output(total_reward, total_steps, total_near_fall,
total_beside_cliff, seed)

# Store results for comparison
results.append({
    'gamma': gamma,
    'iterations': iter,
    'avg_reward': np.mean(total_reward),
    'avg_steps': np.mean(total_steps),
    'avg_near_fall': np.mean(total_near_fall),
    'avg_beside_cliff': np.mean(total_beside_cliff)
})

# Print comparison summary
print("\nComparison Summary:")
print("γ\tIterations\tAvg Reward\tAvg Steps\tAvg Near Fall\tAvg Beside
Cliff")
print("-" * 80)
for r in results:
    print(f"{r['gamma']:.2f}\t{r['iterations']}\t\
\t{r['avg_reward']:.2f}\t{r['avg_steps']:.2f}\t\
\t{r['avg near fall']:.3f}\t{r['avg beside cliff']:.3f}")
```


Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2838.25it/s]

Results for $\gamma=0.1$:

Converged in 8 iterations

Final Value Function $V(s)$:

```
[[ -1.11111111 -1.11111111 -1.11111111 -1.11111111 -1.11111111
  -1.11111111 -1.11111111 -1.11111111 -1.11111109 -1.11111108
  -1.11110647 -1.11106376]
 [ -1.11111111 -1.11111111 -1.11111111 -1.11111111 -1.11111111
  -1.11111111 -1.11111111 -1.11111111 -1.111111092 -1.11110677
  -1.111024 -1.10978535]
 [ -1.11111111 -1.11111111 -1.11111111 -1.11111111 -1.11111111
  -1.11111111 -1.11111111 -1.11111106 -1.11110967 -1.11106815
  -1.10982811 -1.07275122]
 [ -1.11111111 -34.11111111 -67.11111111 -67.11111111 -67.11111111
  -67.11111111 -67.11111111 -67.11111111 -67.11111106 -67.11110968
  -34.07403131 -1.03575837]]
```

Final Policy:

```
[[0 0 0 1 1 1 1 1 1 1 1]
 [0 0 2 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 1]
 [3 0 3 0 0 0 0 0 0 0 1]]
```

==> Output for seed: 17

==> Average reward: -72.75

==> Average steps: 72.75

==> Average near fall: 0.043727942603400585


==> Average beside cliff: 0.32316583697605666

==> Standard deviation of reward: 32.06598665252638

==> Standard deviation of steps: 32.06598665252638

==> Standard deviation of near fall: 0.02882878103475689

==> Standard deviation of beside cliff: 0.11787447173812425



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2932.38it/s]

Results for $\gamma=0.5$:

Converged in 24 iterations

Final Value Function $V(s)$:

```
[[-1.99999998 -1.99999991 -1.99999963 -1.99999852 -1.99999412
  -1.99997704 -1.99991242 -1.99967708 -1.99886794 -1.99633367
  -1.98963153 -1.97763157]
 [-1.99999998 -1.99999991 -1.99999964 -1.99999848 -1.99999356
  -1.99997277 -1.99988501 -1.99951744 -1.99800604 -1.99203682
  -1.97052609 -1.91052627]
 [-1.99999998 -1.99999994 -1.99999973 -1.99999878 -1.99999449
  -1.99997457 -1.99988018 -1.99942151 -1.99713145 -1.98536117
  -1.92299875 -1.58210525]
 [-2.          -34.99999999 -67.99999995 -67.9999998  -67.99999908
  -67.99999576 -67.99998003 -67.99990358 -67.99952191 -67.99756019
  -34.65383312 -1.26368421]]
```

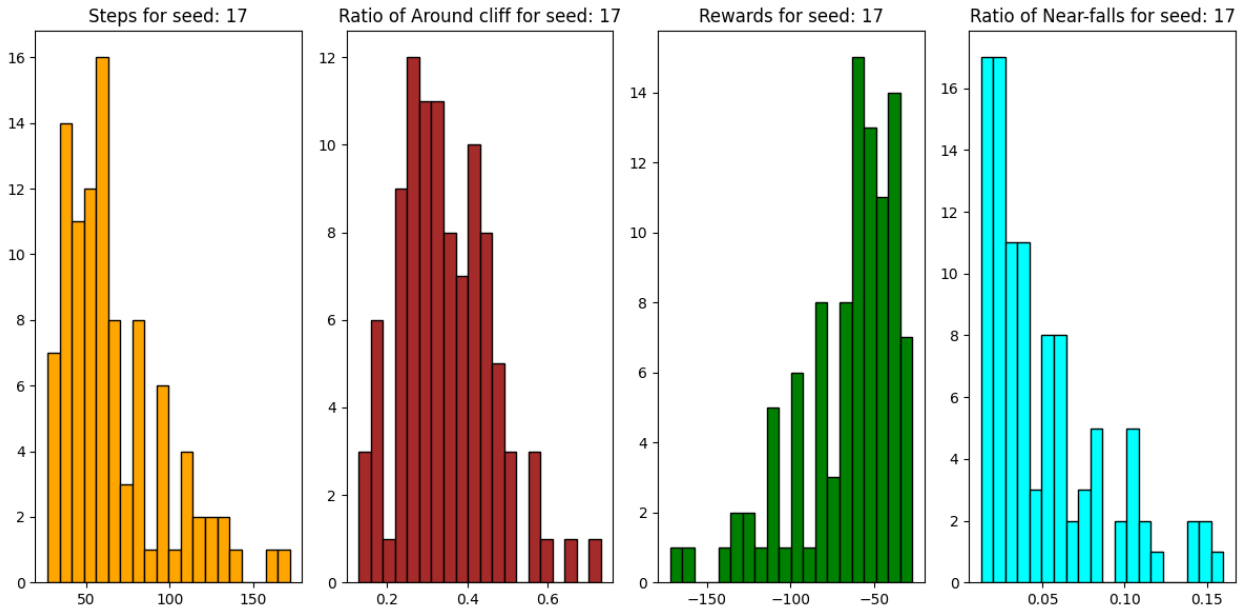
Final Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 3 3 0 0 0 0 0 0 0 1 1]]
==> Output for seed: 17
==> Average reward: -66.42
==> Average steps: 66.42
==> Average near fall: 0.050750757961032476
==> Average beside cliff: 0.3451705764063589
```

```

==> Standard deviation of reward: 30.035705418717903
==> Standard deviation of steps: 30.035705418717903
==> Standard deviation of near fall: 0.03508011006307254
==> Standard deviation of beside cliff: 0.11436303062198966

```



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2580.03it/s]

Results for $\gamma=0.8$:

Converged in 17 iterations

Final Value Function $V(s)$:

```

[[ -4.9986267  -4.99759669  -4.99553287  -4.99168868  -4.98457607
  -4.97152048  -4.94786452  -4.90594181  -4.83449661  -4.72121416
  -4.56594567  -4.4181314 ]
 [ -4.99872172  -4.99785799  -4.99602668  -4.99256779  -4.9860637
  -4.9738168   -4.9506856   -4.90684335  -4.8236515   -4.66739328
  -4.38821919  -3.98172995]
 [ -4.99893303  -4.99834411  -4.99699937  -4.99437684  -4.98934599
  -4.97960692  -4.96036314  -4.92106927  -4.83680325  -4.64329143
  -4.15814634  -2.81153816]
 [ -4.99939032 -37.99923326 -70.99887467 -70.99817533 -70.99683377
 -70.99423668 -70.98910501 -70.97862664 -70.9561557  -70.90455255
 -36.44200978  -1.74974351]]

```

Final Policy:

```


[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 2]]

```

```

[0 0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 0 0 0 0 0 0 0 0 0 0 1]]
==> Output for seed: 17
==> Average reward: -66.38
==> Average steps: 66.38
==> Average near fall: 0.04400681769422257
==> Average beside cliff: 0.36877514413634066
==> Standard deviation of reward: 21.380261925430194
==> Standard deviation of steps: 21.380261925430194
==> Standard deviation of near fall: 0.03183331974095798
==> Standard deviation of beside cliff: 0.10760734004524505

```



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 2787.08it/s]

Results for $\gamma=0.9$:

Converged in 11 iterations

Final Value Function $V(s)$:

[-9.90225461	-9.86967242	-9.81685059	-9.74139885	-9.63470991
	-9.48483183	-9.27613692	-8.99013931	-8.61004928	-8.13507602
	-7.61569565	-7.21649485]			
[-9.90627415	-9.87905154	-9.83125243	-9.76188734	-9.6628246
	-9.52180399	-9.32084683	-9.0336091	-8.62170563	-8.03281507
	-7.22012833	-6.2886598]			
[-9.91522336	-9.89591351	-9.85877008	-9.80206765	-9.71956807
	-9.60033464	-9.42641006	-9.16685207	-8.76282107	-8.08751254
	-6.82940567	-4.12371134]			

```
[ -9.9364178 -42.93062473 -75.9194817 -75.90247098 -75.8777211
-75.84195107 -75.7897737 -75.7119063 -75.590697 -75.38810444
-39.02974704 -2.2371134 ]]
```

Final Policy:

```
[[0 1 1 1 1 1 1 1 1 1 1 1]
[0 1 1 1 1 1 1 1 1 1 1 1]
[0 0 0 0 0 0 0 0 0 0 0 1]
[3 0 0 0 0 0 0 0 0 0 1 1]]
```

==> Output for seed: 17

==> Average reward: -65.96

==> Average steps: 65.96

==> Average near fall: 0.042792288574813646


==> Average beside cliff: 0.35695207277556706

==> Standard deviation of reward: 25.084624772956044

==> Standard deviation of steps: 25.084624772956044

==> Standard deviation of near fall: 0.027349501117388673

==> Standard deviation of beside cliff: 0.10466226021240337



Running the policy for 100 episodes...

100%|██████████| 100/100 [00:00<00:00, 3302.52it/s]

Results for $\gamma=0.99$:

Converged in 4 iterations

Final Value Function $V(s)$:

```
[[ -43.84044516 -42.13863835 -39.83181964 -37.30538546 -
34.61170433
```

```

-31.77867821 -28.82754649 -25.78912346 -22.72795229 -
19.78715325
-17.26742435 -15.71405118]
[ -44.11303591 -42.69208103 -40.53497464 -38.09922915 -
35.46326655
-32.66249691 -29.70922836 -26.60148012 -23.32717409 -
19.87618982
-16.31374977 -13.15993151]
[ -44.72700106 -43.66602555 -41.871904 -39.65325405 -
37.15993952
-34.45905359 -31.56863496 -28.46394518 -25.05396118 -
21.09967216
-15.97794646 -7.97429463]
[ -46.35267936 -79.00255681 -111.4104967 -110.67834221 -
109.85554842
-108.96425606 -108.01041792 -106.98587029 -105.86057557 -
104.55566019
-54.56910652 -3.63151723]]


```

Final Policy:

```

[[0 1 1 1 1 1 1 1 1 1 1 1]
 [0 1 1 1 1 1 1 1 1 1 1 1]
 [0 0 0 0 0 0 0 0 0 0 0 1]
 [3 0 0 0 0 0 0 0 0 0 1 1]]
==> Output for seed: 17
==> Average reward: -65.41
==> Average steps: 65.41
==> Average near fall: 0.048144234296345303
==> Average beside cliff: 0.3345128994561721
==> Standard deviation of reward: 22.546882267843596
==> Standard deviation of steps: 22.546882267843596
==> Standard deviation of near fall: 0.032079142020878804
==> Standard deviation of beside cliff: 0.1011065352027616

```



Comparison Summary:

γ	Iterations	Avg Reward	Avg Steps	Avg Near Fall	Avg Beside Cliff

0.10	8	-72.75	72.75	0.044	0.323
0.50	24	-66.42	66.42	0.051	0.345
0.80	17	-66.38	66.38	0.044	0.369
0.90	11	-65.96	65.96	0.043	0.357
0.99	4	-65.41	65.41	0.048	0.335